

ECE4122/6122 Final Project Report
Three Approaches to Implementing the Two-Dimensional
Fourier Transform using Thread, MPI, and Cuda

Professor: Dr. Brian Swenson

Wootae Song:	wootae@gatech.edu
Alberto Li:	ali97@gatech.edu
Edward Zhou:	edz@gatech.edu
Benjamin Thornbloom:	bthornbloom3@gatech.edu

Due: Dec 13, 2018

Abstract

The goal of the final project was to implement the two-dimensional fourier transform using three different methods: `std::Thread` library, MPI, and Cuda. We have also implemented the inverse discrete fourier transform to validate our results. Each method of implementation was uniquely written for optimized runtime based on the computing model. We compared the performance of each method against each other. This information is shown in the results section.

STD::Thread

The first implementation of the two-dimensional fourier transform was a multi-threaded approach using the `std::thread` library and running on a machine with eight cores. When the source code is compiled, it will produce an executable of the name “p31”. This method was implemented using an N^2 approach as suggested by the project description.


MPI

The second implementation of the two-dimensional fourier transform was an MPI approach running on a coc-ice MPI job using a pbs script with 8 MPI ranks. When the source code is compiled, it will produce an executable of the name “p32”. This method was implemented using the Danielson-Lanczos algorithm, an $N \log(N)$ approach as suggested by the project description.

The third implementation of the two-dimensional fourier transform was a GPU approach using CUDA and running on a coc-ice job queue. When the source code is compiled, it will produce an executable of the name “p33”. This method was implemented using an N^2 approach.

For the Thread computing model, we have implemented the inverse discrete fourier transform.

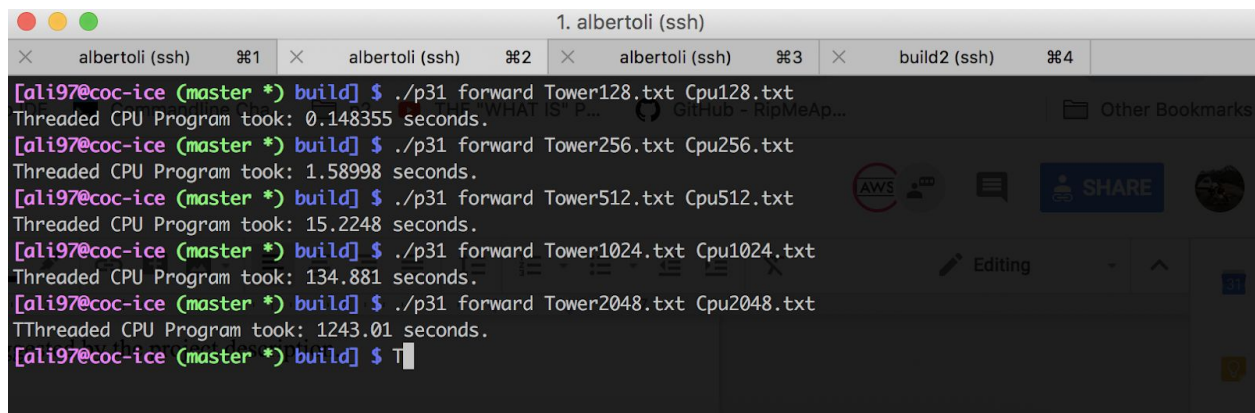
IDFT(DFT(Image))) will reproduce the values of the original image. Below is a screenshot of the first few entries of the IDFT on Tower256.txt.



Results

The `std::Thread`, MPI, and CUDA computing methods were tested with text files of the following sizes: 128x128, 256x256, 512x512, 1024x1024, and 2048x2048. The execution times for each of the methods were timed. `std::Thread` and MPI were timed using `std::chrono`'s high resolution clock at a microsecond granularity and CUDA was timed by submitting pbs scripts and looking at the walltime due to `nvcc`'s bug with `chrono`.

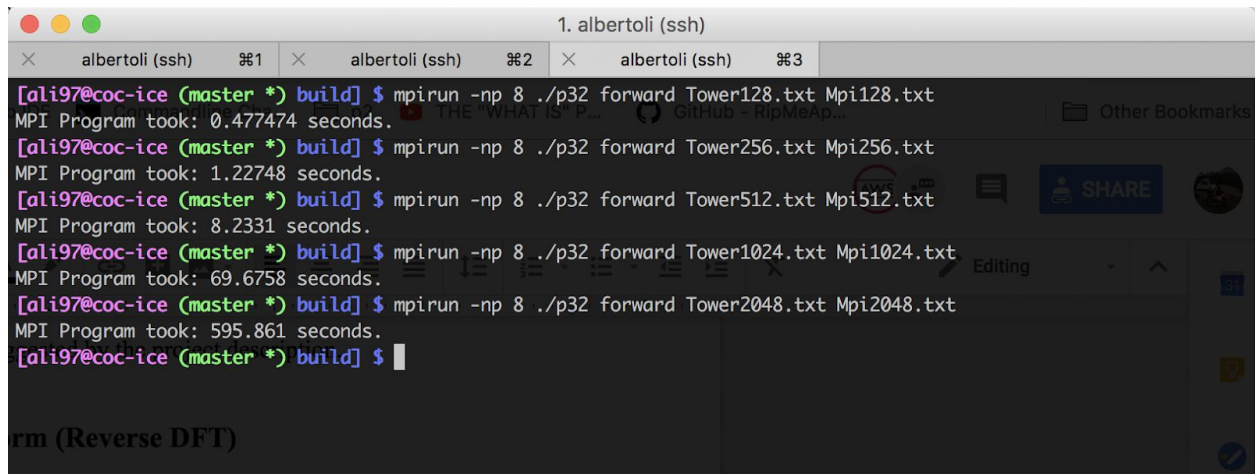
`std::Thread`:



```
1. albertoli (ssh)
× albertoli (ssh) %1 × albertoli (ssh) %2 × albertoli (ssh) %3 × build2 (ssh) %4
[ali97@coc-ice (master *) build] $ ./p31 forward Tower128.txt Cpu128.txt
Threaded CPU Program took: 0.148355 seconds.
[ali97@coc-ice (master *) build] $ ./p31 forward Tower256.txt Cpu256.txt
Threaded CPU Program took: 1.58998 seconds.
[ali97@coc-ice (master *) build] $ ./p31 forward Tower512.txt Cpu512.txt
Threaded CPU Program took: 15.2248 seconds.
[ali97@coc-ice (master *) build] $ ./p31 forward Tower1024.txt Cpu1024.txt
Threaded CPU Program took: 134.881 seconds.
[ali97@coc-ice (master *) build] $ ./p31 forward Tower2048.txt Cpu2048.txt
Threaded CPU Program took: 1243.01 seconds.
[ali97@coc-ice (master *) build] $ T
```

Figure 2. `std::Thread`'s timing for all the sizes

MPI:



```
1. albertoli (ssh)
albertoli (ssh) %1 albertoli (ssh) %2 albertoli (ssh) %3
[ali97@coc-ice (master *) build] $ mpirun -np 8 ./p32 forward Tower128.txt Mpi128.txt
MPI Program took: 0.477474 seconds.
[ali97@coc-ice (master *) build] $ mpirun -np 8 ./p32 forward Tower256.txt Mpi256.txt
MPI Program took: 1.22748 seconds.
[ali97@coc-ice (master *) build] $ mpirun -np 8 ./p32 forward Tower512.txt Mpi512.txt
MPI Program took: 8.2331 seconds.
[ali97@coc-ice (master *) build] $ mpirun -np 8 ./p32 forward Tower1024.txt Mpi1024.txt
MPI Program took: 69.6758 seconds.
[ali97@coc-ice (master *) build] $ mpirun -np 8 ./p32 forward Tower2048.txt Mpi2048.txt
MPI Program took: 595.861 seconds.
[ali97@coc-ice (master *) build] $
```

Figure 3. MPI's timing for all the sizes

CUDA:

PBS JOB 34728.ice-sched.pace.gatech.edu



System Account <adm@ice-sched.pace.gatech.edu>

6:38 PM



받는 사람: a20154920@gmail.com

PBS Job Id: 34728.ice-sched.pace.gatech.edu

Job Name: se2_wootae_song

Exec host: rich133-k33-17.pace.gatech.edu/2

Execution terminated

Exit_status=0

resources_used.cput=00:00:00

resources_used.energy_used=0

resources_used.mem=0kb

resources_used.vmem=0kb

resources_used.walltime=00:00:00

Error_Path: rich133-s30-12.pace.gatech.edu:/nv/coc-ice/ali97/ece6122/assignments/hw5/p3/build/output

Output_Path: rich133-s30-12.pace.gatech.edu:/nv/coc-ice/ali97/ece6122/assignments/hw5/p3/build/output

Figure 4. CUDA's timing for 128x128

PBS JOB 34731.ice-sched.pace.gatech.edu



System Account <adm@ice-sched.pace.gatech.edu>

6:41 PM

받는 사람: a20154920@gmail.com

PBS Job Id: 34731.ice-sched.pace.gatech.edu
Job Name: se2_wootae_song
Exec host: rich133-k33-17.pace.gatech.edu/1
Execution terminated
Exit_status=0
resources_used.cput=00:00:00
resources_used.energy_used=0
resources_used.mem=0kb
resources_used.vmem=0kb
resources_used.walltime=00:00:01
Error_Path: rich133-s30-12.pace.gatech.edu:/nv/coc-ice/ali97/ece6122/assignments/hw5/p3/build/output
Output_Path: rich133-s30-12.pace.gatech.edu:/nv/coc-ice/ali97/ece6122/assignments/hw5/p3/build/output

Figure 5. CUDA's timing for 256x256

PBS JOB 34732.ice-sched.pace.gatech.edu



System Account <adm@ice-sched.pace.gatech.edu>

6:42 PM

받는 사람: a20154920@gmail.com

PBS Job Id: 34732.ice-sched.pace.gatech.edu
Job Name: se2_wootae_song
Exec host: rich133-k33-17.pace.gatech.edu/1
Execution terminated
Exit_status=0
resources_used.cput=00:00:00
resources_used.energy_used=0
resources_used.mem=0kb
resources_used.vmem=0kb
resources_used.walltime=00:00:01
Error_Path: rich133-s30-12.pace.gatech.edu:/nv/coc-ice/ali97/ece6122/assignments/hw5/p3/build/output
Output_Path: rich133-s30-12.pace.gatech.edu:/nv/coc-ice/ali97/ece6122/assignments/hw5/p3/build/output

Figure 6. CUDA's timing for 512x512

PBS JOB 34734.ice-sched.pace.gatech.edu



System Account <adm@ice-sched.pace.gatech.edu>

6:42 PM

받는 사람: a20154920@gmail.com

PBS Job Id: 34734.ice-sched.pace.gatech.edu

Job Name: se2_wootae_song

Exec host: rich133-k33-17.pace.gatech.edu/1

Execution terminated

Exit_status=0

resources_used.cput=00:00:00

resources_used.energy_used=0

resources_used.mem=0kb

resources_used.vmem=0kb

resources_used.walltime=00:00:02

Error_Path: rich133-s30-12.pace.gatech.edu:/nv/coc-ice/ali97/ece6122/assignments/hw5/p3/build/output

Output_Path: rich133-s30-12.pace.gatech.edu:/nv/coc-ice/ali97/ece6122/assignments/hw5/p3/build/output

Figure 7. CUDA's timing for 1024x1024

PBS JOB 34736.ice-sched.pace.gatech.edu



System Account <adm@ice-sched.pace.gatech.edu>

6:44 PM



받는 사람: a20154920@gmail.com

PBS Job Id: 34736.ice-sched.pace.gatech.edu

Job Name: se2_wootae_song

Exec host: rich133-k33-17.pace.gatech.edu/1

Execution terminated

Exit_status=0

resources_used.cput=00:00:00

resources_used.energy_used=0

resources_used.mem=0kb

resources_used.vmem=0kb

resources_used.walltime=00:00:06

Error_Path: rich133-s30-12.pace.gatech.edu:/nv/coc-ice/ali97/ece6122/assignments/hw5/p3/build/output

Output_Path: rich133-s30-12.pace.gatech.edu:/nv/coc-ice/ali97/ece6122/assignments/hw5/p3/build/output

Figure 8. CUDA's timing for 2048x2048

The results of each DFT computing method were compared against the results of the corresponding sizes (128x128, 256x256,...,2048x2048) and the results of each matched. The table below summarizes the results:

	A	B	C	D
1		p31	p32	p33
2	128	0.148355	0.477474	0
3	256	1.58998	1.22748	1
4	512	15.2248	8.2331	1
5	1024	134.881	69.6758	2
6	2048	1243.01	595.861	6

Figure 9. Results matrix in seconds.

As expected, CUDA took significantly less time than the other methods and std::Thread took the most time among all the methods. For the 2048x2048 DFT, CUDA only took 6 seconds to run while std::Thread took around 21 minutes which was 207 times longer. Additionally for the 2048x2048 DFT, CUDA only took 6 seconds to run while MPI took around 10 minutes which was 99 times longer. Lastly for the 2048x2048 DFT, MPI only took 10 minutes to run while CUDA took around 21 minutes which was 2 times longer. As N increases, the difference in execution time of each of the computing methods will increase drastically as std::Thread is $O(n^2)$, MPI is $O(n \log(n))$, and CUDA is $O(1)$.

These performance comparisons are visualized separately (Figures. 10, 11, 12) and together (Figure. 13). By comparing the runtime results of the three different DFT implementations, the following graphs below were produced.




Figure 10. Time (s) vs Size (NxN) for std::Thread




Figure 11. Time (s) vs Size (NxN) for MPI




Figure 12. Time (s) vs Size (NxN) for CUDA




Figure 13. Time (s) vs Size (NxN) for std::Thread, MPI, and CUDA

In conclusion, the results of our three unique implementations based on the programming models: std::Thread, MPI, and CUDA suggest that GPU computation was the most efficient, MPI was the second most efficient, and std::Thread was the least efficient.